

Report for the practical assignment Automated Reasoning 2IMF25

Group 37

Luxue Wen	Kaiwen Tan	Yan Liu
student number: 2271796	student number: 2291797	student number: 2176246
email: l.wen@student.tue.nl	email: k.tan@student.tue.nl	email: y.liu17@student.tue.nl

September, 2025

Problem 1: Magic Factory

1 Modeling the Problem

1.1 Parameters

- $T = 8$: number of trucks.
- $C = 8000$: capacity (kg) per truck.
- $P = 8$: maximum number of pallets per truck.
- Pallet types:
 - Nuzzles: 4 pallets, 700 kg each.
 - Skipples: 8 pallets, 1000 kg each, require cooling.
 - Crottles: 10 pallets, 2500 kg each.
 - Dupples: 20 pallets, 200 kg each.
 - Prittles: unlimited supply, 400 kg each (objective: maximize).
- Cooling: only 3 trucks can carry skipples.

1.2 Decision Variables

Let:

$$x_{i,t} \in \mathbb{N}$$

be the number of pallets of type i assigned to truck t , where $i \in \{\text{nuzzle, prittle, skipple, crottle, duple}\}$ and $t \in \{1, \dots, 8\}$.

Additionally:

$$y_t \in \{0, 1\}$$

indicates whether truck t is equipped with cooling (1) or not (0).

1.3 Constraints

Truck capacity (weight):

$$\forall t: \sum_i (w_i \cdot x_{i,t}) \leq C$$

Truck capacity (pallet count):

$$\forall t : \sum_i x_{i,t} \leq P$$

Supply limits:

$$\sum_t x_{\text{nuzzle},t} = 4, \quad \sum_t x_{\text{skipple},t} = 8, \quad \sum_t x_{\text{crottle},t} = 10, \quad \sum_t x_{\text{dupple},t} = 20$$

Cooling requirement:

$$\sum_t y_t = 3, \quad \forall t : x_{\text{skipple},t} \leq P \cdot y_t$$

Nuzzle distribution:

$$\forall t : x_{\text{nuzzle},t} \leq 1$$

Explosive combination (Part b only):

$$\forall t : x_{\text{prittle},t} \cdot x_{\text{crottle},t} = 0$$

1.4 Objective Function

$$\text{maximize} \quad \sum_t x_{\text{prittle},t}$$

2 Solver Implementation using Z3

- Integer variables n, p, s, c, d represent pallets per truck; boolean variables *cool* indicate cooled trucks.
- Truck-level constraints enforce:
 - Maximum 8 pallets per truck
 - Weight limits
 - Nuzzle distribution (at most 1 per truck)
 - Skipples only on cooled trucks
- Global constraints ensure delivery of all nuzzles, skipples, cattles, dupples, and exactly 3 cooled trucks.
- Part (b) is handled by preventing prittles and cattles in the same truck using `Or(p[i]==0, c[i]==0)`.
- The solver maximizes the total number of prittles.

3 Results Part (a)

3.1 Truck Assignments

Table 1: Truck assignment of pallets for Part (a)

Truck	Nuzzles	Prittles	Skipples	Crottles	Dupples	Total Weight (kg)	Cooling
1	0	1	0	2	5	6400	False
2	1	5	0	0	2	3100	False
3	1	0	0	2	5	6700	False
4	0	0	0	2	6	6200	False
5	1	7	0	0	0	3500	False
6	0	3	1	2	2	7600	True
7	1	1	6	0	0	7100	True
8	0	5	1	2	0	8000	True
Total	4	22	8	10	20	48600	3

3.2 Solver Performance and Optimization Results

Table 2: Optimization results summary (template)

Metric	Value
Total Prittles delivered	22
Solver runtime (s)	0.0241
Solver status	sat

- From the obtained solution, we can see that the solver successfully enforced all constraints. in Part (a), the solver was able to place them flexibly, which allowed the total number of prittles to reach the maximum.
- Observe solver runtime and efficiency. The solver completed the optimization very quickly, taking only 0.0241 s.

4 Results Part (b)

4.1 Truck Assignments

Table 3: Truck assignment of pallets for Part (b)

Truck	Nuzzles	Prittles	Skipples	Crottles	Dupples	Total Weight (kg)	Cooling
1	0	4	4	0	0	5600	True
2	1	0	0	2	5	6700	False
3	1	0	0	2	5	6700	False
4	0	0	2	2	4	7800	True
5	0	8	0	0	0	3200	False
6	1	0	2	2	1	7900	True
7	1	0	0	2	5	6700	False
8	0	8	0	0	0	3200	False
Total	4	20	8	10	20	48400	3

4.2 Solver Performance and Optimization Results

Table 4: Optimization results summary

Metric	Value
Total Prittles delivered	20
Solver runtime (s)	0.5243
Solver status	sat

- In Part (b), the additional constraint preventing prittles and crottles from being assigned to the same truck adds a layer of complexity. The solver correctly respects this restriction along with all previous constraints. As a result, trucks carry either prittles or crottles but not both, and the nuzzles are still properly distributed.
- The runtime increased to 0.5243 s compared to Part (a), reflecting the added combinatorial challenge introduced by the explosive combination constraint. Despite this, the solver still finds the optimal solution reliably.

Problem 2: Poster Printing

1 Modeling the Problem

1.1 Parameters

- $N_{canvas} = 3$: number of each canvas.
- $N_{poster} = 12$: number of each canvas.
- $w = [5, 5, 4, 3, 7, 6, 5, 4, 6, 4, 6, 5]$: width of each poster.
- $h = [6, 6, 10, 11, 7, 10, 13, 10, 9, 15, 10, 10]$: height of each poster.
- $price = [10, 14, 13, 15, 10, 17, 21, 16, 16, 23, 19, 17]$: price of each poster.
- $W = [12, 12, 20]$: width of each canvas.
- $H = [12, 12, 20]$: height of each canvas.
- $cost = [30, 30, 90]$: cost of each canvas.
- $minimal_profit = 60$: minimal profit of printing.

1.2 Decision Variables

To fit posters into canvases, we introduce the following variables:

- $z_{c,p} \in \mathbb{B}$ for $c = 1, \dots, N_{canvas}$, $p = 1, \dots, N_{poster}$: the value of $z_{c,p}$ will be true if and only if post[p] will be printed on canvas[c]
- $r_{c,p} \in \mathbb{B}$ for $c = 1, \dots, N_{canvas}$, $p = 1, \dots, N_{poster}$: the value of $r_{c,p}$ will be true if and only if post[p] will be turned 90°
- $u_c \in \mathbb{B}$ for $c = 1, \dots, N_{canvas}$: the value of u_c will be true if and only if canvas[c] will be used
- $x_{c,p}, y_{c,p} \in \mathbb{N}$ for $c = 1, \dots, N_{canvas}$, $p = 1, \dots, N_{poster}$: the values of $x_{c,p}$ and $y_{c,p}$ indicate the bottom-left coordinate (x, y) of poster[p] placed in canvas[c]
- $w_eff_i, h_eff_i, w_eff_j, h_eff_j \in \mathbb{N}$: the width and height of post[i] and post[j]
- $total_profit \in \mathbb{N}$: the value of total profit after printing

1.3 Constraints

For $post[p]$, it cannot be printed more than once. This is expressed by the formula

$$\sum_i z_{p,i} \leq 1$$

Next we determine that whether $post[p]$ fits into $canvas[c]$. This is expressed by the formula

$$z_{c,p} \Rightarrow \left((\neg r_{c,p} \wedge x_{c,p} \geq 0 \wedge y_{c,p} \geq 0 \wedge x_{c,p} + w[p] \leq W[c] \wedge y_{c,p} + h[p] \leq H[c] \wedge w[p] \leq W[c] \wedge h[p] \leq H[c]) \right. \\ \left. \vee (r_{c,p} \wedge x_{c,p} \geq 0 \wedge y_{c,p} \geq 0 \wedge x_{c,p} + w[p] \leq W[c] \wedge y_{c,p} + h[p] \leq H[c]) \wedge h[p] \leq W[c] \wedge w[p] \leq H[c] \right)$$

Additionally, every two posters $post[i]$ and $post[j]$ should have no overlap. This is expressed by the formula

$$(z_{c,i} \wedge z_{c,j}) \Rightarrow \left((x_{c,i} + w_{eff_i} \leq x_{c,j}) \vee (x_{c,j} + w_{eff_j} \leq x_{c,i}) \right. \\ \left. \vee (y_{c,i} + h_{eff_i} \leq y_{c,j}) \vee (y_{c,j} + h_{eff_j} \leq y_{c,i}) \right)$$

Then, we associate $canvase[c]$ and $poster[p]$. This is expressed by the formula

$$z_{c,p} \Rightarrow u_c$$

Finally, we set the minimal profit. This is expressed by the formula

$$total_profit \geq mininal_profit$$

1.4 Calculation Function

The calculation function of w_{eff_i} , h_{eff_i} , w_{eff_j} , h_{eff_j} is expressed by the formula

$$w_{eff_i} = \begin{cases} h_i, & \text{if } r_{c,i} = 1 \\ w_i, & \text{if } r_{c,i} = 0 \end{cases} \quad h_{eff_i} = \begin{cases} w_i, & \text{if } r_{c,i} = 1 \\ h_i, & \text{if } r_{c,i} = 0 \end{cases} \\ w_{eff_j} = \begin{cases} h_j, & \text{if } r_{c,j} = 1 \\ w_j, & \text{if } r_{c,j} = 0 \end{cases} \quad h_{eff_j} = \begin{cases} w_j, & \text{if } r_{c,j} = 1 \\ h_j, & \text{if } r_{c,j} = 0 \end{cases}$$

The calculation function of $total_profit$ is expressed by the formula

$$\sum_{\substack{c=1, \dots, N_{canvas} \\ p=1, \dots, N_{poster} \\ z_{c,p}=true}} price_p - \sum_{\substack{c=1, \dots, N_{canvas} \\ u_c=true}} cost_c$$

2 Solver Implementation using Z3

- Define variables for each poster's placement, rotation, and which canvas it's assigned to, along with flags indicating whether a canvas is used.
- Make sure each poster is placed at most once, fits within the chosen canvas, and can be rotated if needed.
- Prevent posters from overlapping on the same canvas by considering their positions and effective sizes.
- Automatically exclude posters that are too big for any canvas.
- Link canvas usage with poster assignment, so a canvas is only counted if it actually holds posters.
- Calculate total profit as poster revenues minus canvas costs, and enforce that it meets the minimal profit requirement.
- Use the Z3 solver to find a feasible arrangement that satisfies all these constraints and maximizes the overall profit.

3 Results Part

3.1 posters Assignments with three canvases

Table 5: posters Assignments with three canvases for Part (a)

canvas	p0	p1	p2	p3	p4	p5	p6	p7	p8	p9	p10	p11	price	cost	profit
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1	10	14	0	15	0	0	0	16	0	0	0	0	55	30	
2	0	0	13	0	0	17	21	0	16	23	19	17	126	90	
Total													181	120	61

It is possible to obtain profit at least 60.

3.2 posters Assignments with two small canvases

Table 6: posters Assignments with three canvases for Part (a)

canvas	p0	p1	p2	p3	p4	p5	p6	p7	p8	p9	p10	p11	price	cost	profit
0	10	14	0	0	0	0	0	0	0	0	0	0	19	43	30
1	0	0	0	15	0	0	0	16	0	0	0	17	48	30	
Total													91	60	31

The highest profit created by the two small canvases is 31.

3.3 Solver Performance

- In Part (a), the run time is around 900 ms.
- In Part (a), the run time is around 2 s.

Problem 3: Dinner Organisation

1 Modeling the Problem

1.1 Parameters

- $N = 10$: number of participants.
- $H = 5$: number of houses, each occupied by one couple.
- $K = 5$: number of courses.
- Each house can host at most 5 participants per course.
- Each couple must host exactly 2 courses.
- Every course is served in two houses simultaneously.
- Guest distribution: in each course, each hosting couple has exactly 3 guests.

1.2 Decision Variables

- Integer variable

$$\text{attend}_{p,k} \in \{0, 1, 2, 3, 4\}, \quad p = 0, \dots, 9, \quad k = 0, \dots, 4,$$

indicating the house where participant p attends course k .

- Boolean variable

$$\text{serve}_{h,k} \in \{\text{False}, \text{True}\}, \quad h = 0, \dots, 4, \quad k = 0, \dots, 4,$$

indicating whether house h hosts course k .

- Auxiliary integer variable

$$\text{meet}_{p,q} \in \{0, \dots, 5\}, \quad 0 \leq p < q \leq 9,$$

counting the number of courses where participants p and q attend the same house:

$$\text{meet}_{p,q} = \sum_{k=0}^4 [\text{attend}_{p,k} = \text{attend}_{q,k}].$$

1.3 Constraints

Course hosting.

$$\forall k : \sum_{h=0}^4 \text{If}(\text{serve}_{h,k}, 1, 0) = 2, \quad \forall h : \sum_{k=0}^4 \text{If}(\text{serve}_{h,k}, 1, 0) = 2$$

Couples attend their own house when hosting. For house h , let the couple be participants $2h$ and $2h + 1$:

$$\forall h, k : \text{serve}_{h,k} \Rightarrow (\text{attend}_{2h,k} = h \wedge \text{attend}_{2h+1,k} = h)$$

Non-hosting participants attend exactly one house per course.

$$\forall p, k : \sum_{h=0}^4 [\text{attend}_{p,k} = h] = 1$$

Each serving house has exactly 5 participants (couple + 3 guests).

$$\forall h, k : \sum_{p=0}^9 [\text{attend}_{p,k} = h] = \text{If}(\text{serve}_{h,k}, 5, 0)$$

Meeting counts. For each pair $p \neq q$:

$$\text{meet}_{p,q} = \sum_{k=0}^4 [\text{attend}_{p,k} = \text{attend}_{q,k}]$$

Meeting frequency constraints.

- property 1: $\text{meet}_{p,q} \geq 1$
- property 2: $\text{meet}_{p,q} \leq 3$

Couples never meet outside their own house (property 3).

$$\forall h, k : (\text{attend}_{2h,k} = \text{attend}_{2h+1,k}) \Rightarrow (\text{attend}_{2h,k} = h)$$

Distinct guests per house (property 4). For house h hosting courses k_1 and k_2 , the guest sets must be disjoint:

$$\forall h, i \notin \{2h, 2h + 1\} : \sum_{t \in \{k_1, k_2\}} [\text{attend}_{i,t} = h] \leq 1$$

2 Solver Implementation using Z3

- Integer variables $\text{attend}_{p,k}$ encode participant attendance at houses during courses.
- Boolean variables $\text{serve}_{h,k}$ encode which houses host each course.
- Couples are fixed in their own house during hosting.
- Attendance constraints ensure each serving house has exactly 5 participants (couple + 3 guests).
- Each couple hosts exactly 2 courses, and each course is served in exactly 2 houses.
- Part (a):
 - Require each two people meet at least once.
 - Require couples never meet outside or guests are distinctive, but not both.
- Part (b):
 - Require each pair meets at most 3 times.
 - Couples never meet outside.
 - Guests are distinctive per house.

3 Results Part (a)

3.1 Schedule

Table 7: Dinner schedule for Part (a) - Scenario 1 (Property 1 + Property 3)

Course	Hosting Houses	Guest Distribution
0	1, 2	1: [0, 2, 3, 6, 9], 2: [1, 4, 5, 7, 8]
1	0, 1	0: [0, 1, 5, 7, 9], 1: [2, 3, 4, 6, 8]
2	0, 2	0: [0, 1, 2, 7, 8], 2: [3, 4, 5, 6, 9]
3	3, 4	3: [0, 2, 4, 6, 7], 4: [1, 3, 5, 8, 9]
4	3, 4	3: [1, 3, 4, 6, 7], 4: [0, 2, 5, 8, 9]

Table 8: Dinner schedule for Part (a) - Scenario 2 (Property 1 + Property 4)

Course	Hosting Houses	Guest Distribution
0	2, 3	2: [1, 4, 5, 8, 9], 3: [0, 2, 3, 6, 7]
1	0, 4	0: [0, 1, 2, 5, 7], 4: [3, 4, 6, 8, 9]
2	2, 4	2: [3, 4, 5, 6, 7], 4: [0, 1, 2, 8, 9]
3	1, 3	1: [0, 1, 2, 3, 4], 3: [5, 6, 7, 8, 9]
4	0, 1	0: [0, 1, 4, 6, 8], 1: [2, 3, 5, 7, 9]

Table 9: Dinner schedule for Part (a) - Scenario 3 (Property 1 + Property 3 + Property 4)

Course	Hosting Houses	Guest Distribution
No results		

3.2 Solver Performance and Optimization Results

Table 10: Optimization results summary - Scenario 1

Metric	Value
Feasible schedule found	Yes
Solver runtime (s)	0.0405
Solver status	sat

Table 11: Optimization results summary - Scenario 2

Metric	Value
Feasible schedule found	Yes
Solver runtime (s)	0.0189
Solver status	sat

Table 12: Optimization results summary - Scenario 3

Metric	Value
Feasible schedule found	No
Solver runtime (s)	4.7659
Solver status	unsat

- Scenario 1 confirms that a schedule satisfying Property 1 and Property 3 exists.
- Scenario 2 confirms that a schedule satisfying Property 1 and Property 4 exists.
- Scenario 3 confirms that satisfying Properties 1, 3, and 4 simultaneously is impossible (expected UNSAT).

4 Results Part (b)

4.1 Schedule

Table 13: Dinner schedule for Part (b) (template: fill with solver output)

Course	Hosting Houses	Guest Distribution
0	1, 4	1: [1, 2, 3, 4, 6], 4: [0, 5, 7, 8, 9]
1	1, 4	1: [0, 2, 3, 5, 7], 4: [1, 4, 6, 8, 9]
2	2, 3	2: [0, 2, 4, 5, 8], 3: [1, 3, 6, 7, 9]
3	0, 3	0: [0, 1, 3, 4, 9], 3: [2, 5, 6, 7, 8]
4	0, 2	0: [0, 1, 2, 7, 8], 2: [3, 4, 5, 6, 9]

4.2 Solver Performance and Optimization Results

Table 14: Optimization results summary (Part b)

Metric	Value
Feasible schedule found	Yes
Solver runtime (s)	0.0180
Solver status	sat

- The solver verified that property 2 can indeed be satisfied together with both 3 and 4.

Problem 4: Program Verification

Problem statement.

We are given the following C program:

```
uint32_t x = nondet();
int y = 1;
while (x < (int)1e9) {
    x = x + y;
    y = 2*y + 1;
}
```

Tasks: (a) find the smallest initial value of x_0 such that the loop executes exactly 15 times; (b) analyze whether overflow or underflow can occur.

Symbolic Encoding of Program Semantics

To verify this program, we represent its semantics symbolically as a system of constraints.

State variables. Each loop iteration i maintains two symbolic variables x_i and y_i .

$$\begin{aligned} x_{i+1} &= x_i + y_i, \\ y_{i+1} &= 2y_i + 1, \end{aligned}$$

starting from $y_0 = 1$ and an unconstrained x_0 (the nondeterministic input).

Loop condition. The loop continues as long as $x < 10^9$. For “exactly 15 iterations”, this condition translates to:

$$x_{14} < 10^9, \quad x_{15} \geq 10^9.$$

These inequalities precisely encode the boundary between iteration 14 and 15.

Z3 Constraint Model Construction

We create the following constraints in Z3:

- Declare integer arrays $x[0..15]$, $y[0..15]$.
- Add recurrence constraints:

$$y[i+1] = 2*y[i] + 1, \quad x[i+1] = x[i] + y[i].$$

- Add the initial condition: $y[0] = 1$.

- Add the loop boundary condition:

$$x[14] < 1e9, \quad x[15] \geq 1e9.$$

- Use Z3's `Optimize()` object to minimize `x[0]`.

The resulting Z3 script is:

```
from z3 import *

N = 15
x = [Int(f"x{i}") for i in range(N+1)]
y = [Int(f"y{i}") for i in range(N+1)]

opt = Optimize()
opt.add(y[0] == 1)
for i in range(N):
    opt.add(y[i+1] == 2*y[i] + 1)
    opt.add(x[i+1] == x[i] + y[i])
opt.add(x[14] < 10**9, x[15] >= 10**9)
opt.minimize(x[0])
opt.check()
print(opt.model())
```

Z3 Solving Process

Z3 expands all 15 iterations, substitutes constraints recursively, and performs *constraint propagation* until all variables are expressed in terms of x_0 . Internally, it builds a system of linear arithmetic constraints of the form:

$$x_{14} = x_0 + 32752, \quad x_{15} = x_0 + 65519,$$

which is the same result obtained by manual symbolic expansion.

Then the optimizer solves the inequalities:

$$x_{14} < 10^9, \quad x_{15} \geq 10^9,$$

and determines the minimal feasible x_0 satisfying both, yielding:

$$x_0^{min} = 999\,934\,481.$$

Thus, the solver reproduces automatically what we would obtain through algebraic reasoning, but via pure logical constraint solving.

Result Interpretation

For $x_0 = 999,934,481$:

$$x_{14} = 999,967,233 < 10^9, \quad x_{15} = 1,000,000,000 \geq 10^9.$$

Therefore, the loop executes **exactly 15 iterations**. This matches both analytical reasoning and Z3's output.

(b) Overflow Analysis

To verify the absence of integer overflow, we modeled the program using **Z3 BitVec (32-bit)** variables, which exactly simulate C arithmetic. Each addition operates modulo 2^{32} , so any overflow would automatically wrap around to zero and can be detected by a constraint of the form `x[i+1] < x[i]`.

k	$y_k = 2^{k+1} - 1$	$\Delta x = y_k$	x_k (with $x_0 = 0$)
0	1	+1	1
1	3	+3	4
2	7	+7	11
\vdots	\vdots	\vdots	\vdots
28	536,870,911	+536,870,911	536,870,880
29	1,073,741,823	+1,073,741,823	1,073,741,793 ($\geq 1e9$)

Table 15: Growth of x and y during the loop.

Z3 model.

```

x = BitVec('x', 32)
y = BitVec('y', 32)
# same recurrence as before, unrolled 15 steps
x_next = x + y
y_next = 2*y + 1
# overflow condition: x_next < x (wrap-around)

```

The solver checks whether such a model can exist. The result is **unsat**, which means no overflow occurs in 15 iterations under 32-bit semantics.

Interpretation. In unsigned arithmetic, x is monotonically increasing and cannot decrease unless a wrap-around occurs. Since the solver found no model with $x_{i+1} < x_i$, no wrap-around is possible. For the signed variable y , its value stays positive and below $2^{31} - 1$ within 15 iterations, so no signed overflow occurs either.

Why underflow is not modeled. Both variables are monotonically increasing: $y_{k+1} = 2y_k + 1 > y_k$ and $x_{k+1} = x_k + y_k > x_k$. Negative or decreasing values can never occur, so underflow is semantically impossible. Only overflow needs to be checked.

Z3 therefore proves that the program is safe from both signed and unsigned overflow within its execution bounds.

Overflow Visualization (15 iterations)

To complement the solver-based proof, the following table shows how x and y grow numerically over the 15 iterations. All values remain far below the 32-bit limits ($2^{31} - 1$ for signed, $2^{32} - 1$ for unsigned).

k	$y_k = 2^{k+1} - 1$	x_k (with $x_0 = 0$)	Comparison to limits
13	16,383	16,366	$\ll 2^{31} - 1$ (INT_MAX)
14	32,767	32,736	$\ll 2^{31} - 1$
15	65,535	65,471	$\ll 2^{31} - 1, \ll 2^{32} - 1$

Table 16: Variable magnitudes within 15 iterations (actual program case).

After 15 iterations, both x and y remain far below the 32-bit bounds. Z3 confirms symbolically that no overflow or wrap-around can occur, and the numerical visualization provides an intuitive confirmation.

Runtime Discussion

- (a) Integer minimization (15 steps): 0.007s in Z3.
- (b) BitVec 32-bit model (15 steps): 0.009s in Z3; even 29 steps remain below 0.01s.

Conclusions

- The loop executes exactly 15 times for $x_0 = 999\,934\,481$.
- Neither overflow nor underflow can occur.
- Analytical and Z3 results match perfectly.

Problem 5: Configurable Systems Testingg

1 Modeling and Implementation

1.1 Methodology

This report details the analysis of five configurable software systems, whose valid configurations are defined by constraints in DIMACS CNF format. The core of the methodology was to use a Binary Decision Diagram (BDD) to create a compact and canonical representation of each system’s valid configuration space. All subsequent analyses were performed on this BDD representation using the `Oxidd` library in Python.

The primary objectives for each system were to:

- (i) Determine the size of the BDD in terms of node count.
- (ii) Calculate the total number of valid configurations, $|V|$.
- (iii) Perform Uniform Random Sampling (URS) to find the selection ratio of a specific feature (x_{42}).
- (iv) Count the total number of valid pairwise feature interactions.
- (v) Find the size of a small test suite, $|B|$, that achieves pairwise coverage.

1.2 Implementation Details

For each system, the DIMACS file was parsed to extract the clauses and the recommended variable order from ‘c vo’ lines. This order was applied during BDD construction to minimize its size. A weighted random walk algorithm was implemented for the URS task, and a greedy set-cover algorithm was used for the pairwise cover generation.

2 Results and Analysis

The implemented solution was run on the five provided systems. The results reveal significant differences in their structural complexity and computational demands.

2.1 buildroot.dimacs

Table 17: Results for `buildroot.dimacs`

Metric	Value
(i) BDD Nodes	1,000
(ii) Valid Configurations $ V $	$\approx 1.98 \times 10^{158}$
(iii) URS Ratio k_1/k_0 for x_{42}	5024 / 4976
(iv) Pairwise Interactions	621,270
(v) Cover Set Size $ B $	Incomplete (96% finished)

2.2 toybox.dimacs

Table 18: Results for `toybox.dimacs`

Metric	Value
(i) BDD Nodes	949
(ii) Valid Configurations $ V $	$\approx 1.45 \times 10^{17}$
(iii) URS Ratio k_1/k_0 for x_{42}	0 / 10000
(iv) Pairwise Interactions	256,494
(v) Cover Set Size $ B $	Incomplete (65% finished)

2.3 busybox.dimacs

Table 19: Results for `busybox.dimacs`

Metric	Value
(i) BDD Nodes	3,076
(ii) Valid Configurations $ V $	$\approx 3.42 \times 10^{194}$
(iii) URS Ratio k_1/k_0 for x_{42}	1644 / 8356
(iv) Pairwise Interactions	1,322,443
(v) Cover Set Size $ B $	Incomplete (projected >161 hours)

2.4 embtoolkit.dimacs

Table 20: Results for `embtoolkit.dimacs`

Metric	Value
(i) BDD Nodes	172,157
(ii) Valid Configurations $ V $	(An integer with approx. 340 digits)
(iii) URS Ratio k_1/k_0 for x_{42}	Not Run (projected >179 hours)
(iv) Pairwise Interactions	Not Run
(v) Cover Set Size $ B $	Not Run

2.5 uClinux.dimacs

Table 21: Results for `uClinux.dimacs`

Metric	Value
(i) BDD Nodes	2,667
(ii) Valid Configurations $ V $	$\approx 1.63 \times 10^{91}$
(iii) URS Ratio k_1/k_0 for x_{42}	0 / 10000
(iv) Pairwise Interactions	3,013,528
(v) Cover Set Size $ B $	Incomplete (projected >10 hours)

3 Conclusion on Time and Machine Limitations

The BDD-based approach proved effective for representing the configuration spaces. However, the results clearly demonstrate significant computational limitations when analyzing systems with high structural complexity. For systems like `busybox`, `embtoolkit`, and `uClinux`, the sheer size of the BDD or the number of interactions made some analysis tasks (especially tasks (iii) and (v)) computationally intractable.

The incomplete and “Not Run” results are a direct consequence of these time and machine constraints, as the algorithm’s runtime grows substantially with the problem’s complexity, making it infeasible to obtain a complete result for the most complex systems within a reasonable timeframe on the available hardware.

Problem 6: Finite State Automata

Problem statement.

This problem aims to analyze a finite-state automaton (FSA) using Binary Decision Diagrams (BDDs). The specific tasks are as follows:

- (a) Compactly represent the FSA by a BDD using a suitable encoding and variable order.
- (b) Determine whether the accepting states are reachable from the initial states by:
 - (i) taking only 0-transitions,
 - (ii) taking only 1-transitions, or
 - (iii) alternating 0–1 transitions, i.e., by a path $s_0 \xrightarrow{0} s_1 \xrightarrow{1} s_2 \xrightarrow{0} \dots s_k$ with $s_k \in F$.

Justify the approach and clearly state what each BDD represents.

The experiment is demonstrated using the input file `bakery.1.c.ba`.

Methodology: BDD-based Reachability Analysis

A Binary Decision Diagram (BDD) compactly encodes Boolean relations and enables efficient symbolic traversal of large state spaces. Each system state is represented as a vector of Boolean variables.

The analysis process (implemented in `finite_state_automata02.py`) proceeds as follows:

1. Parse and read the automaton file to extract the number of states, transitions, initial and accepting states;
2. Compute the number of bits required for binary encoding of states;
3. Define the BDD variable set $\{x_0, \dots, x_{m-1}, y_0, \dots, y_{m-1}, a\}$: x_i/y_i encode the i -th bit of current/next state in an interleaved order $(x_0, y_0, x_1, y_1, \dots)$, and a encodes the 0/1 action label;
4. Build the transition relation: For every edge (s, α, t) in the automaton, the program encodes the source state s , the destination state t , and the transition label α as Boolean vectors. Each transition is represented in the BDD as a conjunction of:
 - the bits of the current state (x_i variables),
 - the action bit a indicating whether $\alpha = 0$ or 1 , and
 - the bits of the next state (y_i variables).

All transitions are then combined with a disjunction (OR) to form the global transition relation T .

5. Use an interleaved variable order (e.g., $x_0, y_0, x_1, y_1, \dots$) to reduce node count and cache conflicts;
6. Compute reachable states iteratively as $R_{k+1} = R_k \cup \text{Image}(R_k, T)$ until a fixpoint is reached.

For the three restricted modes, the program constructs filtered transition relations T_0, T_1, T_{alt} corresponding to 0-only, 1-only, and alternating 0–1 transitions. We start from $S_0 = I, S_1 = \perp$; the first expansion uses $a = 1$ -edges (to S_1), then $a = 0$ -edges (back to S_0), iterating to a fixpoint.

Implementation Overview

The Python script `finite_state_automata02.py` executes the following steps:

1. Parse the automaton file (`bakery.1.c.ba`);
2. Build the symbolic BDD-based transition relation T ;
3. Perform reachability analysis for the three transition restrictions;
4. Cross-check the symbolic results with an explicit state-space traversal;
5. Print runtime statistics and analysis logs.

System Parameters

Parameter	Value
Input file	<code>bakery.1.c.ba</code>
States	1506
Transitions	2697
State range	0–1505
State encoding bits	11
BDD variables	interleaved: $x_0, y_0, \dots, x_{10}, y_{10}$, then a

Table 22: System parameters for BDD-based FSA analysis.

Experimental Results

The program successfully parsed the automaton file `bakery.1.c.ba`, containing 1506 states and 2697 transitions. Each state was encoded using 11 bits with interleaved variables $(x_0, \dots, x_{10}, y_0, \dots, y_{10})$, plus the action variable a . The BDD transition relation was built successfully.

Reachability analyses produced the following results:

- (i) Only 0-transitions: accepting states reachable (True)
- (ii) Only 1-transitions: accepting states reachable (True)
- (iii) Alternating 0–1 transitions: accepting states reachable (True)

Explicit cross-check confirmed the symbolic results:

- 0-edges = 2338, 1-edges = 359
- reachable states = $\{1, 16, (\neg S_0 \rightarrow 3, \neg S_1 \rightarrow 2)\}$
- all verdicts = True

All transitive closures converged in one iteration, and the total runtime was approximately 0.003 seconds.

Interpretation

If the automaton does not explicitly define accepting states, all states are considered accepting by default. Therefore, any reachable state is automatically an accepting state. This explains why all three transition restrictions yield True results.

Symbolic vs. Explicit Comparison

Explicit analysis is only used to validate symbolic results. BDD-based symbolic reasoning performs Boolean operations directly over state encodings, achieving the same accuracy without enumerating all transitions.

Method	Feature	Description
Explicit traversal	Enumerates all states	Intuitive to implement but computationally expensive for large systems
Symbolic (BDD) analysis	Operates on Boolean functions	Compact representation, high efficiency, consistent with explicit results

Table 23: Comparison between explicit traversal and symbolic BDD analysis.

Conclusions

- Under all three transition restrictions (0-only, 1-only, alternating 0–1), accepting states are reachable.
- Symbolic and explicit analyses produce identical results.
- The automaton is highly connected across different transition modes.
- BDD-based reachability converges in one iteration and executes efficiently (runtime ≈ 0.003 s).

Appendix Table

The following table summarizes experimental results for the analyzed automaton files. Each entry reports the total number of states and transitions, the reachability results for 0-only, 1-only, and alternating 0–1 transitions, and the total runtime in seconds.

File name	States	Transitions	0-trans	1-trans	0–1 alternating	Runtime (s)
bakery.1.c.ba	1506	2697	True	True	True	0.03
bakery.2.c.ba	1146	2085	True	True	True	0.02
fischer.2.c.ba	21733	67590	True	True	True	1.25
fischer.3.1.c.ba	638	1401	True	True	True	0.02
fischer.3.2.c.ba	1536	3856	True	True	True	0.05
fischer.3.c.ba	637	1400	True	True	True	0.02
mcs.1.2.c.ba	7968	21509	True	True	True	0.30
NFA_hard.1.ba	4959	26850	True	True	True	0.34
NFA_hard.2.ba	4857	24147	True	True	True	0.35
phils.1.1.c.ba	161	466	True	True	True	0.01
phils.2.c.ba	581	2350	True	True	True	0.02

Table 24: Summary of FSA analysis results for multiple input files.

Problem 7: Multi-Machine Task Scheduling

Problem Statement

Eight machines are waiting in a busy factory, and there are 20 tasks that must be processed. Each task has a certain processing time, power consumption, and a deadline, and some tasks depend on others. All task parameters are provided in the file `task_data.json`.

- Each machine can only work on one task at a time.
 - Tasks with dependencies are processed in the correct order.
 - The total power consumption of all running tasks does not exceed the factory’s power limit.
 - Each task should ideally finish by its deadline; otherwise, it counts as tardy.
- (a) Find a feasible task schedule that minimizes the total completion time and the tasks’ tardiness.
(b) Give a possible optimization to speed up the process. Justify your answers.

1 Modeling the Problem

1.1 Parameters

- $num_tasks = 20$: number of tasks.
- $num_machines = 8$: number of machines.
- $Pmax$: power constraint.
- $duration$: duration time of each task.
- $power$: power consumption of each task.
- $deadline$: deadline of each task.
- $precedences$: list of task order constraints.

1.2 Decision Variables

- $start[i] \in \mathbb{N}$: the start time of task[i].
- $machine[i] \in \mathbb{N}$: task[i] is assigned to machine[i].
- $tardiness[i] \in \mathbb{N}$: the tardiness time of task[i].
- $Cmax \in \mathbb{N}$: the total time to finish all tasks.

1.3 Constraints

Each task must have a non-negative start time,

$$start_i \geq 0, \quad i \in \{0, \dots, num_tasks - 1\}$$

Each task is assigned to a valid machine,

$$0 \leq machine_i < num_machines, \quad i \in \{0, \dots, num_tasks - 1\}$$

Each task's tardiness must be non-negative,

$$tardiness_i \geq 0, \quad i \in \{0, \dots, num_tasks - 1\}$$

For every precedence pair (i, j), task i must finish before task j starts,

$$start_i + duration_i \leq start_j, \quad (i, j) \in precedences$$

For any two tasks i and j assigned to the same machine do not overlap in time,

$$machine_i \neq machine_j \vee start_i + duration_i \leq start_j \vee start_j + duration_j \leq start_i, \quad i \neq j$$

The sum of powers of running tasks must not exceed Pmax,

$$\sum_{i \in \{0, \dots, num_tasks - 1\}} power_i \leq P_{max},$$

Tardiness is calculated as,

$$tardiness_i \geq start_i + duration_i - deadline_i,$$

Cmax needs to be at least as large as all task finish times,

$$C_{max} \geq start_i + duration_i, \quad i \in \{0, \dots, num_tasks - 1\}$$

2 Optimization Strategy

- Before optimization,
 - (a) synchronous power constraints and scheduling constraints,
 - (b) the sum of the power consumptions of all running tasks does not exceed the system limit P_{\max} .
- Optimization strategy:
 - (a) First minimizing the maximum completion time,

$$C_{\max}^* = \min C_{\max}$$

Then, minimizing total tardiness given the fixed completion time,

$$\min \sum_{i=1}^N t_i \text{ s.t. } C_{\max} = C_{\max}^*$$

- (b) pairwise mutual exclusion constraint,

$$(power_i + power_j > P_{\max}) \Rightarrow (start_i + duration_i \leq start_j \vee start_j + duration_j \leq start_i), i < j$$

3 Solver Implementation using Z3

- Define start times, machine assignments, task tardiness, and maximum completion time.
- Add domain constraints, precedence constraints, machine conflict constraints, power constraints, tardiness constraints and maximum completion time constraint to z3 solver.
- Adopt two-phase optimization. First minimize makespan, then minimize total tardiness while keeping makespan fixed.

4 Results Part

4.1 Visualize the results




Figure 1: Task scheduling Gantt chart (before optimization).




Figure 2: Task scheduling Gantt chart (after optimization).

4.2 Solver Performance

- Initially, the total runtime is 47.462405 seconds.
- After optimization, the total runtime is 1.649142 seconds.
- In conclusion, the optimized version saves 96.52% of time.